



EEC1509 - Machine Learning

Lesson #4 - Linear Regression One Variable

Ivanovitch Silva
September, 2018





- We are going to start by covering linear regression
 - One variable
 - Multiples variables
- We discuss the application of linear regression to **housing price prediction**
- Present the notion of a **cost function**
- Introduce the **gradient descent** method for learning.
- Refresher on **linear algebra concepts**.

Update repository

```
git clone https://github.com/ivanovitchm/EEC1509_MachineLearning.git
```

Ou

```
git pull
```



Regression




5 kilograms

200 kilograms

1.5 kilograms



????



1.5 kilograms



????

```

1 import pandas as pd
2 from sklearn import linear_model
3 import matplotlib.pyplot as plt
4
5 #read data
6 df = pd.read_fwf('brain_body.txt')
7 X = df[['Brain']]
8 y = df[['Body']]
9
10
11 #train model on data
12 model = linear_model.LinearRegression()
13 model.fit(X, y)
14
15 #visualize results
16 plt.scatter(X.values,y.values)
17 plt.plot(X.values,model.predict(X),color='red')
18 plt.xlabel('Brain weight (kg)')
19 plt.ylabel('Body weight (kg)')
20
21 plt.show()

```


1_hello_world.ipynb (15min)



Training Process

Training Data

features					labels
x1	x2	x3	xn	y

Select Features



Find the best model that best **approximates** training set.

Training is the most computationally intensive step.

$O(n)^2$



$\hat{y} = 2.5x_1 + 11x_2 + 3.5x_3 + \dots + a_n x_n$

Testing Process

Test Data

features					labels
x1	x2	x3	xn	y



Predict label for an instance by plugging features into trained model.

$\hat{y} = 2.5x_1 + 11x_2 + 3.5x_3 + \dots + a_n x_n$

Predicting is computationally cheap.




features					labels	predicted labels	
x1	x2	x3	xn	y		\hat{y}

Compute Error Metric

$O(1)$

Linear Regression - Housing Price



Lot Area	Overall Qual	Year Built	Yr Sold	SalePrice
31770	6	1960	2010	215000
11622	5	1961	2010	105000
14267	6	1958	2010	172000
11160	7	1968	2010	244000
13830	5	1997	2010	189900

One variable
Multiple variables



Reproduce this fig. using plotly and scikit-learn



Read sections 1 and 2 (15min)
2. The linear Regression Model.ipynb



Linear Regression with One Variable

Notation:

- m - number of training examples
- X 's - input variable/features
- y 's - output variable/ target variable

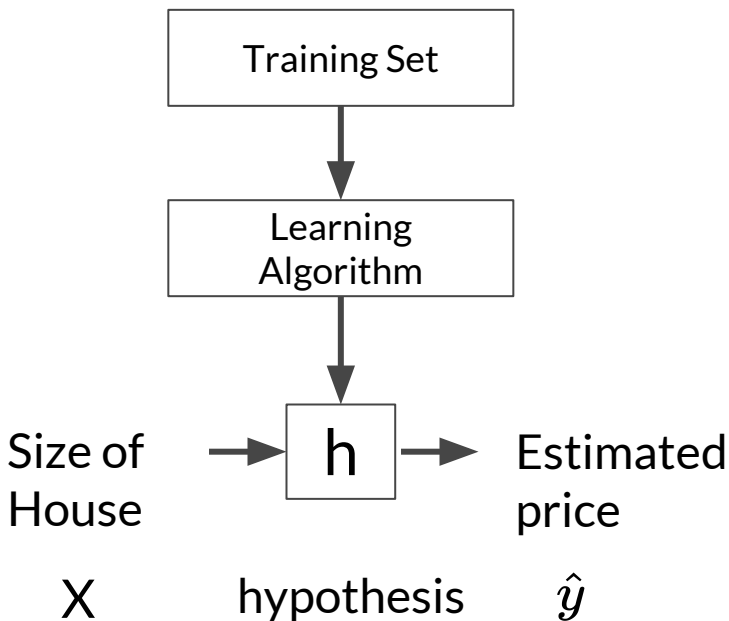
$$\begin{array}{ll} X^{(1)} = 31770 & y^{(1)} = 215000 \\ X^{(2)} = 11622 & y^{(2)} = 105000 \\ X^{(3)} = 14267 & y^{(3)} = 172000 \end{array}$$

$(X^{(i)}, y^{(i)}) = i^{\text{th}}$ training example

$m = 1465$


<div>X</div> <div>y</div>	
Lot Area SalePrice	
31770	215000
11622	105000
14267	172000
11160	244000
13830	189900

Model Representation (linear reg. one variable)



How do we represent h ?

$$h_{\theta}(x) = \theta_0 + \theta_1 x$$



Cost Function

$$f(x)$$

"minimize the error"

Cost Function (Linear Reg. One Var.)

Hypothesis $h_{\theta}(x) = \theta_0 + \theta_1 x$

θ_i = parameters

How to choose θ_i ?

$m = 1465$

Training Set


X	y
Lot Area	SalePrice
31770	215000
11622	105000
14267	172000
11160	244000
13830	189900

Cost Function


Intuition #01

(linear reg. One var)


1. Let's plot:



1. Here are our **points** and **prediction**
line



1. Prediction line $Y = 0.180 X$ has an error value of 2,058. Let's try other lines.



Cost Function (square error function)

Training Set (m instances)

	Lot Area	SalePrice	
$x^{(1)}$	31770	215000	$y^{(1)}$
$x^{(2)}$	11622	105000	$y^{(2)}$
$x^{(3)}$	14267	172000	$y^{(3)}$
$x^{(4)}$	11160	244000	$y^{(4)}$
$x^{(5)}$	13830	189900	$y^{(5)}$

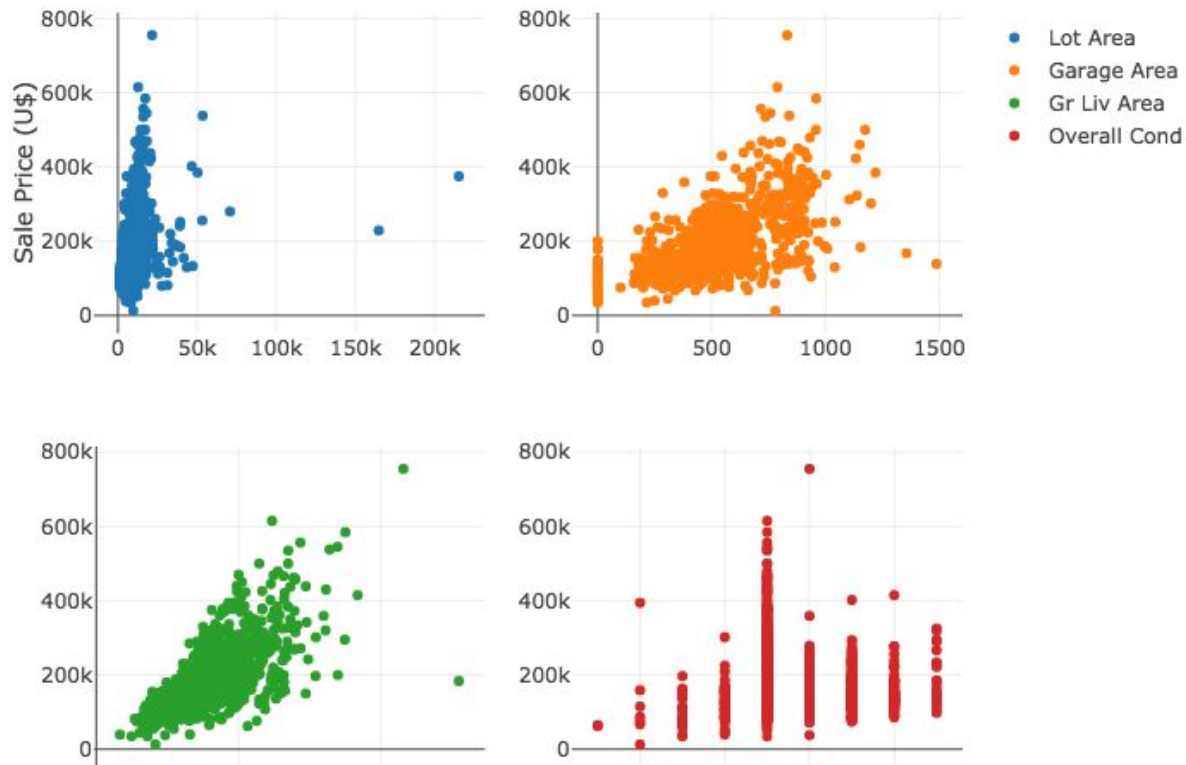
$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m [h_{\theta}(x^{(i)}) - y^{(i)}]^2$$

$$h_{\theta}(x) = \theta_0 + \theta_1 x$$

Idea:

- choose θ_0, θ_1 so that $h_{\theta}(x)$ is close to y for our training examples $(x^{(i)}, y^{(i)})$
- minimize (θ_0, θ_1)

Cost Function[step #01] - Select the feature x

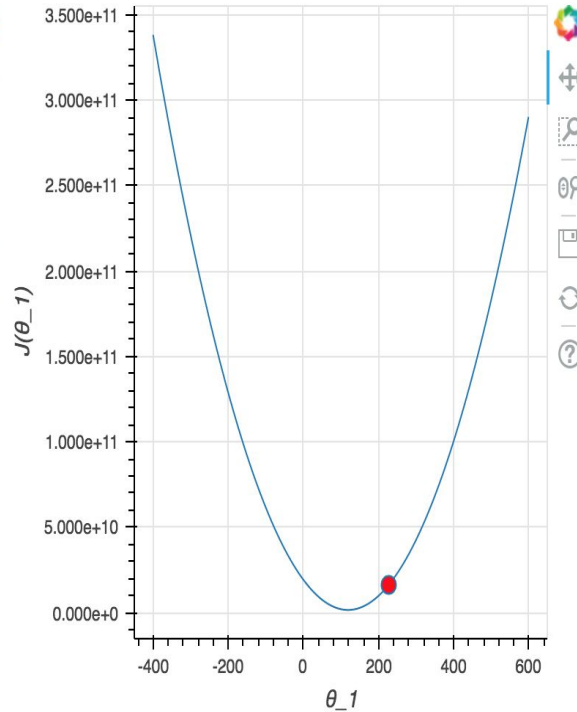



Cost Function[step #01] - Select the feature x_1

```
1 train[['Garage Area', 'Gr Liv Area', 'Overall Cond', 'Lot Area', 'SalePrice']].corr()
```

	Garage Area	Gr Liv Area	Overall Cond	Lot Area	SalePrice
Garage Area	1.000000	0.473506	-0.145705	0.213122	0.625335
Gr Liv Area	0.473506	1.000000	-0.134157	0.248676	0.706364
Overall Cond	-0.145705	-0.134157	1.000000	-0.042415	-0.108979
Lot Area	0.213122	0.248676	-0.042415	1.000000	0.267714
SalePrice	0.625335	0.706364	-0.108979	0.267714	1.000000

$J(227) = 47472029566758$



Cost Function

Intuition #01

$(\theta_0 = 0)$

$$\hat{y} = h_{\theta}(x) = \theta_0 + \theta_1 x$$

$$\hat{y} = h_{\theta}(x) = \theta_1 x$$

$$J(\theta_1) = \frac{1}{2m} \sum_{i=1}^m [\theta_1 x^{(i)} - y^{(i)}]^2$$

A1: 227





Cost Function

Intuition #02

(linear reg. One var)

Contour



Surface

Cost Function


Intuition #02

(θ_0 and θ_1 are defined)

$$\hat{y} = h_{\theta}(x) = \theta_0 + \theta_1 x$$

$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m \left[h_{\theta}(x^{(i)}) - y^{(i)} \right]^2$$

Gradient Descent (linear reg. One var)






Algorithm - Idea

Have some function $J(\theta_0, \theta_1)$

Want $\min_{\theta_0, \theta_1} J(\theta_0, \theta_1)$

Start with some θ_0, θ_1
Keep changing θ_0, θ_1 to reduce $J(\theta_0, \theta_1)$
until we hopefully end up at a minimum





repeat until converge {

$$\theta_j = \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1)$$

}

α - learning rate

repeat until converge {

$$\theta_j = \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1)$$

}

Correct update

$$aux_0 = \theta_0 - \alpha \frac{\partial}{\partial \theta_0} J(\theta_0, \theta_1)$$

$$aux_1 = \theta_1 - \alpha \frac{\partial}{\partial \theta_1} J(\theta_0, \theta_1)$$

$$\theta_0 = aux_0$$

$$\theta_1 = aux_1$$

Incorrect update

$$aux_0 = \theta_0 - \alpha \frac{\partial}{\partial \theta_0} J(\theta_0, \theta_1)$$

$$\theta_0 = aux_0$$

$$aux_1 = \theta_1 - \alpha \frac{\partial}{\partial \theta_1} J(\theta_0, \theta_1)$$

$$\theta_1 = aux_1$$

repeat until converge {



$$aux_0 = \theta_0 - \alpha \frac{\partial}{\partial \theta_0} J(\theta_0, \theta_1) = \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m [h_{\theta}(x^{(i)}) - y^{(i)}]$$

$$aux_1 = \theta_1 - \alpha \frac{\partial}{\partial \theta_1} J(\theta_0, \theta_1) = \theta_1 - \alpha \frac{1}{m} \sum_{i=1}^m [h_{\theta}(x^{(i)}) - y^{(i)}] x^{(i)}$$


$$\theta_0 = aux_0$$

$$\theta_1 = aux_1$$


}



Learning rate
tradeoff



Gradient Descent Pitfalls



cost function & gradient descent from linear algebra perspective

Hypothesis

$$\hat{y} = h_{\theta}(x) = \theta_0 + \theta_1 x$$

Gr Liv Area	SalePrice
2480	205000
1829	237000
2673	249000
1005	133500
1768	224900

[Export to plot.ly »](#)

$$\text{hypothesis} = \begin{bmatrix} 1 & 2480 \\ 1 & 1829 \\ 1 & 2679 \\ 1 & 1005 \\ 1 & 1768 \end{bmatrix} \times \begin{bmatrix} \theta_0 \\ \theta_1 \end{bmatrix} = \begin{bmatrix} 2480 \theta_1 + \theta_0 \\ 1829 \theta_1 + \theta_0 \\ 2679 \theta_1 + \theta_0 \\ 1005 \theta_1 + \theta_0 \\ 1768 \theta_1 + \theta_0 \end{bmatrix}$$

```
def cost_function(X, y, theta):
    return np.sum(np.square(np.matmul(X, theta) - y)) / (2 * len(y))
```

$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m [h_{\theta}(x^{(i)}) - y^{(i)}]^2$$

Gr Liv Area	SalePrice
2480	205000
1829	237000
2673	249000
1005	133500
1768	224900

[Export to plot.ly »](#)

$$J(\theta_0, \theta_1) = \frac{1}{2 \times 5} \sum \left(\begin{bmatrix} 2480 \theta_1 + \theta_0 \\ 1829 \theta_1 + \theta_0 \\ 2679 \theta_1 + \theta_0 \\ 1005 \theta_1 + \theta_0 \\ 1768 \theta_1 + \theta_0 \end{bmatrix} - \begin{bmatrix} 205000 \\ 237000 \\ 249000 \\ 133500 \\ 224900 \end{bmatrix} \right)^2$$


```
def gradient_descent(X, y, alpha, iterations, theta):
    m = len(y)
    all_thetas = [theta]

    for i in range(iterations):
        t0 = theta[0] - (alpha / m) * np.sum(np.dot(X, theta) - y)
        t1 = theta[1] - (alpha / m) * np.sum((np.dot(X, theta) - y) * X[:,1])
        theta = np.array([t0, t1])
        all_thetas.append([t0,t1])

    return theta, np.array(all_thetas)
```

repeat until converge {

- Involves calculations over the full training set X at each gradient step
- It uses the whole batch of training data at every step.
- This is why the algorithm called **Batch Gradient Descent**

$$aux_0 = \theta_0 - \alpha \frac{\partial}{\partial \theta_0} J(\theta_0, \theta_1) = \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m [h_{\theta}(x^{(i)}) - y^{(i)}]$$

$$aux_1 = \theta_1 - \alpha \frac{\partial}{\partial \theta_1} J(\theta_0, \theta_1) = \theta_1 - \alpha \frac{1}{m} \sum_{i=1}^m [h_{\theta}(x^{(i)}) - y^{(i)}] x^{(i)}$$

$$\theta_0 = aux_0$$

$$\theta_1 = aux_1$$

}

Stochastic Gradient Descent

Mini-batch Gradient Descent

